

# (e)BPF, perf, tracing

28.02.2019  
Maciej Borzęcki  
CEHUG Łódź

# This talk is about

- Berkeley Packet Filter
- extended BPF
- SECCOMP
- Tracing
- Having fun while reading code

BPF

# BPF design goals

- It must be protocol independent. The kernel should not have to be modified to add new protocol support
- It must be general. The instruction set should be rich enough to handle unforeseen uses
- Packet data references should be minimized
- Decoding an instruction should consist of a single C switch statement
- The abstract machine registers should reside in physical registers

McCanne, Steven; Jacobson, Van (1992-12-19). "The BSD Packet Filter: A New Architecture for User-level Packet Capture" (PDF).

# Demo - tcpdump

# BPF

```
$ tcpdump -d 'ip and tcp port 80'
(000) ldh      [12]
(001) jeq      #0x800          jt 2 jf 12  <-- EtherType 0x0800 IPv4
(002) ldb      [23]
(003) jeq      #0x6           jt 4 jf 12  <-- IPv4 Protocol - TCP
(004) ldh      [20]
(005) jset     #0x1fff         jt 12    jf 6
(006) ldx      4*([14]&0xf)    <-- Internet Header Length
(007) ldh      [x + 14]        <-- TCP source port
(008) jeq      #0x50          jt 11    jf 9
(009) ldh      [x + 16]        <-- TCP destination port
(010) jeq      #0x50          jt 11    jf 12
(011) ret      #262144         <-- accept 256k of packet data
(012) ret      #0             <-- ignore
```

# BPF

```
$ tcpdump -d 'ip and tcp port 80'
```

```
(000) ldh      [12]
```

```
(001) jeq      #0x800          jt 2 jf 12  <-- EtherType 0x0800 IPv4
```

```
(002) ldb      [23]
```

```
(003) jeq      #0x6           jt 4 jf 12  <-- IPv4 Protocol - TCP
```

```
(004) ldh      [20]
```

```
(005) jset     #0x1fff        jt 12 jf 6
```

```
(006) ldx      4*([14]&0xf)
```

```
(007) ldh      [x + 14]
```

```
(008) jeq      #0x50          jt 12 jf 12
```

```
(009) ldh      [x + 16]
```

```
(010) jeq      #0x50          jt 12 jf 12
```

```
(011) ret      #262144
```

```
(012) ret      #0
```

802.3 Ethernet packet and frame structure

Layer	Preamble	Start of frame delimiter	MAC destination	MAC source	802.1Q tag (optional)	Ethertype (Ethernet II) or length (IEEE 802.3)	Payload
	7 octets	1 octet	6 octets	6 octets	(4 octets)	2 octets	46-1500 octets
Layer 2 Ethernet frame	← 64-1522 octets →						

# BPF

```
$ tcpdump -d 'ip and tcp port 80'
(000) ldh      [12]
(001) jeq      #0x800          jt 2 jf 12  <-- EtherType 0x0800 IPv4
(002) ldb      [23]
(003) jeq      #0x6           jt 4 jf 12  <-- IPv4 Protocol - TCP
(004) ldh      [20]
(005) jset     #0x1fff        jt 12 jf 6
(006) ldx      4*([14]&0xf)    <-- Internet Header Length
(007) ldh      [x + 14]       <-- TCP source port
(008) jeq      #0x50
(009) ldh      [x + 16]
(010) jeq      #0x50
(011) ret      #262144
(012) ret      #0
```

IPv4 Header Format																																			
Offsets	Octet	0						1								2								3											
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31		
0	0	Version				IHL				DSCP				ECN				Total Length																	
4	32	Identification																Flags				Fragment Offset													
8	64	Time To Live								Protocol								Header Checksum																	
12	96	Source IP Address																																	
16	128	Destination IP Address																																	
20	160	Options (if IHL > 5)																																	
24	192																																		
28	224																																		
32	256																																		



# BPF

```
$ tcpdump -d 'ip and tcp port 80'
```

```
(000) ldh      [12]
```

```
(001) jeq      #0x800          jt 2 jf 12  <-- EtherType 0x0800 IPv4
```

```
(002) ldb      [23]
```

```
(003) jeq      #0x6           jt 4 jf 12  <-- IPv4 Protocol - TCP
```

```
(004) ldh      [20]
```

```
(005) jset     #0x1fff         jt 12    jf 6
```

```
(006) ldx      4*([14]&0xf)    <-- Internet Header Length
```

```
(007) ldh      [x + 14]        <-- TCP source port
```

```
(008) jeq      #0x50
```

```
(009) ldh      [x + 16]
```

```
(010) jeq      #0x50
```

```
(011) ret      #262144
```

```
(012) ret      #0
```

IPv4 Header Format																																	
Offsets	Octet	0				1								2								3											
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
0	0	Version				IHL				DSCP								ECN				Total Length											
4	32	Identification																Flags				Fragment Offset											
8	64	Time To Live								Protocol								Header Checksum															
12	96	Source IP Address																															
16	128	Destination IP Address																															
20	160	Options (if IHL > 5)																															
24	192																																
28	224																																
32	256																																

# BPF

```
$ tcpdump -d 'ip and tcp port 80'
```

```
(000) ldh      [12]
(001) jeq      #0x800          jt 2 jf 12  <-- EtherType 0x0800 IPv4
(002) ldb      [23]
(003) jeq      #0x6           jt 4 jf 12  <-- IPv4 Protocol - TCP
(004) ldh      [20]
(005) jset     #0x1fff        jt 12   jf 6
(006) ldx      4*([14]&0xf)    <-- Internet Header Length
(007) ldh      [x + 14]        <-- TCP source port
(008) jeq      #0x50          jt 11   jf 9
(009) ldh      [x + 16]        <-- TCP destination port
(010) jeq      #0x50          jt 11   jf 12
(011) ret      #262144
(012) ret      #0
```

TCP Header																																	
Offsets	Octet	0								1								2								3							
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
0	0	Source port																Destination port															
4	32	Sequence number																															
8	64	Acknowledgment number (if ACK set)																															
12	96	Data offset				Reserved 0 0 0			N S	C W R	E C R	U R G	A C K	P C S	R C S	S S Y	F I N	Window Size															

# BPF abstract machine

- 2 32bit registers, A (accumulator), X (index)
- 16 elements stack, each 32bit
- Basic instruction set
  - load/store (ld\*, ld\*, st\*, stx\*)
  - ALU on A, X or K (constant)
  - Jump (conditional, unconditional)
  - opcode:16b jt:8b jf:8b k:32b
- Forward jumps only

# BPF

```
$ tcpdump -d 'ip6 and tcp port 80'
(000) ldh      [12]
(001) jeq      #0x86dd      jt 2 jf 9  <-- EtherType 0x86dd IPv6
(002) ldb      [20]
(003) jeq      #0x6         jt 4 jf 9  <-- IPv6 Protocol - TCP
(004) ldh      [54]
(005) jeq      #0x50        jt 8 jf 6  <-- TCP source port
(006) ldh      [56]
(007) jeq      #0x50        jt 8 jf 9  <-- TCP destination port
(008) ret      #262144
(009) ret      #0
```

# BPF use case - SECCOMP

# SECCOMP

**seccomp** (short for **secure computing mode**) is a computer security facility in the [Linux kernel](#). seccomp was first devised by Andrea Arcangeli in January 2005 for use in public [grid computing](#) and was originally intended as a means of safely running untrusted compute-bound programs. It was merged into the [Linux kernel mainline](#) in kernel version 2.6.12, which was released on March 8, 2005.<sup>[1]</sup> seccomp allows a [process](#) to make a one-way transition into a "secure" state where it cannot make any [system calls](#) except `exit()`, `sigreturn()`, `read()` and `write()` to already-open [file descriptors](#). Should it attempt any other system calls, the kernel will terminate the process with [SIGKILL](#) or [SIGSYS](#)<sup>[2][3]</sup>. In this sense, it does not [virtualize](#) the system's resources but isolates the process from them entirely.

seccomp mode is enabled via the `prctl(2)` [system call](#) using the `PR_SET_SECCOMP` argument, or (since Linux kernel 3.17<sup>[4]</sup>) via the `seccomp(2)` [system call](#).<sup>[5]</sup> seccomp mode used to be enabled by writing to a file, `/proc/self/seccomp`, but this method was removed in favor of `prctl()`.<sup>[6]</sup> In some kernel versions, seccomp disables the [RDTSC](#) [x86](#) instruction, which returns the number of elapsed processor cycles since power-on, used for high-precision timing.<sup>[7]</sup>

# seccomp

- seccomp() system call
- Control which system calls can be executed
- One way street
- Once enabled cannot be dropped
- Preserved across fork() and clone()
- Preserved across execve()
- Filters are BPF programs

# seccomp

- seccomp() system call
- Control which system calls can be executed
- One way street
- Once enabled cannot be dropped
- Preserved across fork() and clone()
- Preserved across execve()
- Filters are BPF programs





# seccomp - libseccomp & Go

```
// default action - ALLOW
filter, err := seccomp.NewFilter(seccomp.ActAllow)

// find syscall number - read()
sysNum, err := seccomp.GetSyscallFromName("read")

// first arg == 3
cond, err := seccomp.MakeCondition(0, seccomp.CompareEqual, 3)

// read(3, ..) = -EPERM
err = filter.AddRuleConditional(sysNum, seccomp.ActErrno.SetReturnCode(1),
    []seccomp.ScmpCondition{cond})
```

# seccomp - BPF program pseudocode

```
# filter for arch x86_64 (3221225534 = 0xc000003e)
if ($arch == 3221225534)
    # filter for syscall "read" (0) [priority: 65533]
    if ($syscall == 0)
        if ($a0.hi32 == 0)
            if ($a0.lo32 == 3)
                action ERRNO(1);
    # default action
    action ALLOW;
# invalid architecture action
action KILL;
```

# seccomp - BPF program disassembly

```
line  OP   JT   JF   K
=====
0000: 0x20 0x00 0x00 0x00000004  ld  $data[4]
0001: 0x15 0x00 0x0a 0xc000003e  jeq 3221225534 true:0002 false:0012
0002: 0x20 0x00 0x00 0x00000000  ld  $data[0]
0003: 0x35 0x00 0x01 0x40000000  jge 1073741824 true:0004 false:0005
0004: 0x15 0x00 0x07 0xffffffff  jeq 4294967295 true:0005 false:0012
0005: 0x15 0x00 0x04 0x00000000  jeq 0      true:0006 false:0010
0006: 0x20 0x00 0x00 0x00000014  ld  $data[20]
0007: 0x15 0x00 0x02 0x00000000  jeq 0      true:0008 false:0010
0008: 0x20 0x00 0x00 0x00000010  ld  $data[16]
0009: 0x15 0x01 0x00 0x00000003  jeq 3      true:0011 false:0010
0010: 0x06 0x00 0x00 0x7fff0000  ret ALLOW
0011: 0x06 0x00 0x00 0x00050001  ret ERRNO(1)
0012: 0x06 0x00 0x00 0x00000000  ret KILL
```

# Demo

# seccomp - BPF demo

```
int main(int argc, char *argv[]) {
    ...
    char *bpf_prog = read_prog(argv[1], &bpf_prog_size);

    prctl(PR_SET_NO_NEW_PRIVS, 1);
    struct sock_fprog prog = {
        .len = bpf_prog_size / sizeof(struct sock_filter),
        .filter = (struct sock_filter *)bpf_prog,
    };
    seccomp(SECCOMP_SET_MODE_FILTER, SECCOMP_FILTER_FLAG_LOG, &prog);
    /* filter attached */
    int fd = open("/proc/self/exe", 0);
    assert(fd == 3);
    char buf;
    /* THIS SHOULD FAIL */
    ssize_t rd = read(fd, &buf, 1);
    assert(rd == -1);
    perror("read() failed");
    ...
}
```

# seccomp & BPF - kernel side

- kernel/seccomp.c: `seccomp_prepare_filter()` ← when setting up the filter
- net/core/filter.c: `bpf_check_classic()`
- kernel/seccomp.c: `seccomp_run_filters()` ← when the filter is applied
- `BPF_PROG_RUN()`, returns seccomp action

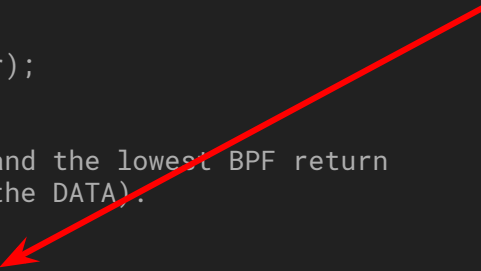
# seccomp & BPF - kernel side

```
static u32 seccomp_run_filters(const struct seccomp_data *sd,
                              struct seccomp_filter **match)
{
    u32 ret = SECCOMP_RET_ALLOW;
    struct seccomp_filter *f =
        READ_ONCE(current->seccomp.filter);

    ...
    /*
     * All filters in the list are evaluated and the lowest BPF return
     * value always takes priority (ignoring the DATA).
     */
    for (; f; f = f->prev) {
        u32 cur_ret = BPF_PROG_RUN(f->prog, sd);

        if (ACTION_ONLY(cur_ret) < ACTION_ONLY(ret)) {
            ret = cur_ret;
            *match = f;
        }
    }
    return ret;
}

struct seccomp_data {
    int nr;
    __u32 arch;
    __u64 instruction_pointer;
    __u64 args[6];
};
```



# seccomp - BPF program disassembly

```
line  OP   JT   JF   K
=====
0000: 0x20 0x00 0x00 0x00000004  ld  $data[4]
0001: 0x15 0x00 0x0a 0xc000003e  jeq 3221225534 true:0002 false:0012
0002: 0x20 0x00 0x00 0x00000000  ld  $data[0]
0003: 0x35 0x00 0x01 0x40000000  jge 1073741824 true:0004 false:0005
0004: 0x15 0x00 0x07 0xffffffff  jeq 4294967295 true:0005 false:0012
0005: 0x15 0x00 0x04 0x00000000  jeq 0      true:0006 false:0010
0006: 0x20 0x00 0x00 0x00000014  ld  $data[20]
0007: 0x15 0x00 0x02 0x00000000  jeq 0      true:0008 false:0010
0008: 0x20 0x00 0x00 0x00000010  ld  $data[16]
0009: 0x15 0x01 0x00 0x00000003  jeq 3      true:0011 false:0010
0010: 0x06 0x00 0x00 0x7fff0000  ret ALLOW
0011: 0x06 0x00 0x00 0x00050001  ret ERRNO(1)
0012: 0x06 0x00 0x00 0x00000000  ret KILL

struct seccomp_data {
    int nr;
    __u32 arch;
    __u64 instruction_pointer;
    __u64 args[6];
};
```



eBPF

# eBPF

- 64bit registers
- r0 - r10
- C compatible ABI
  - Parameters passed in r1 - r5 (r1 carries call context)
  - r6 - r9 callee saved
  - return value in r0
  - r10 frame pointer
- No overhead calling to/from C code
- JITed
- Kernel provides helpers
- Verifier

# eBPF - kernel helper

```
/* prototype for BPF verifier */
const struct bpf_func_proto bpf_get_prandom_u32_proto = {
    .func          = bpf_user_rnd_u32,
    .gpl_only      = false,
    .ret_type      = RET_INTEGER,
};
/* for the actual call from eBPF */
BPF_CALL_0(bpf_user_rnd_u32)
{
    /* Should someone ever have the rather unwise idea to use some
     * of the registers passed into this function, then note that
     * this function is called from native eBPF and classic-to-eBPF
     * transformations. Register assignments from both sides are
     * different, f.e. classic always sets fn(ctx, A, X) here.
     */
    struct rnd_state *state;
    u32 res;
    state = &get_cpu_var(bpf_user_rnd_state);
    res = prandom_u32_state(state);
    put_cpu_var(bpf_user_rnd_state);
    return res;
}
```

# eBPF in the kernel

- cgroups
  - Firewall
  - Device (former device cgroup)
- tc - traffic control
  - classifier
- xtables
- tracing, events
- More peculiar use cases
  - PPP
  - ISDN
  - LIRC

# Demo - tc-bpf

# tc-bpf

```
#include <linux/bpf.h>

#ifndef __section
# define __section(x)  __attribute__((section(x), used))
#endif

__section("classifier") int cls_main(struct __sk_buff *skb)
{
    return -1;
}

char __license[] __section("license") = "GPL";
```

```
$ clang -O2 -emit-llvm -c hello.c -o - | llc -march=bpf -filetype=obj -o hello.o
$ llvm-objdump -S -no-show-raw-insn hello.o
```

# Demo - bpftool

# systemd cgroup packet filter

- Installed for cgroup by systemd
- When service IP address white/black list is used

```
[Unit]
Description=Journal Service
...
[Service]
ExecStart=/usr/lib/systemd/systemd-journald
...
SystemCallFilter=@system-service      ← SECCOMP
SystemCallErrorNumber=EPERM
SystemCallArchitectures=native
...
IPAddressDeny=any                     ← cgroup packet filter
```



# systemd cgroup packet filter

```
$ bpftool prog list
```

```
$ bpftool cgroup tree
```

```
$ bpftool prog dump xlated id <id>
```

```
$ bpftool prog dump jited id <id>
```

# systemd cgroup packet filter - kernel side

```
int __cgroup_bpf_run_filter_skb(struct sock *sk,
                               struct sk_buff *skb,
                               enum bpf_attach_type type)
{
    int ret;

    if (!sk || !sk_fullsock(sk))
        return 0;
    if (sk->sk_family != AF_INET && sk->sk_family != AF_INET6)
        return 0;

    ...
    /* compute pointers for the bpf prog */
    bpf_compute_and_save_data_end(skb, &saved_data_end);
    ret = BPF_PROG_RUN_ARRAY(cgrp->bpf.effective[type], skb,
                            bpf_prog_run_save_cb);
    bpf_restore_data_end(skb, saved_data_end);
    __skb_pull(skb, offset);
    skb->sk = save_sk;
    return ret == 1 ? 0 : -EPERM;
}
```

# Tracing

# Tracing

- functrace
- Trace events
- Kprobes
- Uprobes
- USDT

# Tracing subsystem

- `/sys/debug/kernel/tracing`
- Privileged access only
- See [Documentation/trace/](#)

# perf

- Swiss Army knife of all tracing/profiling
- Profiling

\$ perf stat

\$ perf record/report

- **Tracing**

\$ perf trace

\$ perf record -e ...

\$ perf probe

# Demo - SECCOMP

# perf trace

```
$ sudo perf trace ./demo prog.bpf
? (      ): demo/17292 ... [continued]: execve()) = 0
2.480 ( 0.139 ms): demo/17292 brk(      ) = 0x1d39000
...
6.878 ( 0.038 ms): demo/17292 seccomp(op: FILTER, flags: 0x2, uargs: 0x7ffd734af360) = 0
12.306 ( 0.017 ms): demo/17292 openat(dfd: CWD, filename: 0x4020c0) = 3
12.306 ( 0.042 ms): demo/17292 ... [continued]: read()) = -1 EPERM Operation not permitted
...
```



# kprobe - perf probe & trace

```
$ sudo perf probe -v seccomp_run_filters
```

```
Added new event:
```

```
  probe:seccomp_run_filters (on seccomp_run_filters)
```

You can now use it in all perf tools, such as:

```
  perf record -e probe:seccomp_run_filters -aR sleep 1
```

```
$ sudo perf trace ./demo prog.bpf
```

```
? (      ): demo/17292 ... [continued]: execve()) = 0
```

```
2.480 ( 0.139 ms): demo/17292 brk(      ) = 0x1d39000
```

```
...
```

```
6.878 ( 0.038 ms): demo/17292 seccomp(op: FILTER, flags: 0x2, uargs: 0x7ffd734af360) = 0
```

```
12.306 ( 0.017 ms): demo/17292 openat(dfd: CWD, filename: 0x4020c0) = 3
```

```
12.306 ( 0.042 ms): demo/17292 ... [continued]: read()) = -1 EPERM Operation not permitted
```

```
...
```

# kretprobe - perf probe & trace

```
$ sudo perf probe -v 'seccomp_run_filters%return' '$retval:x32'
```

You can now use it in all perf tools, such as:

```
perf record -e probe:seccomp_run_filters__return -aR sleep 1
```

```
$ sudo perf trace ./demo prog.bpf
```

```
0.000 demo/17473 syscalls:sys_enter_read:fd: 3</home/guest/seccomp/prog.bpf>, buf: 0xcc4a80, ..
```

```
1.004 demo/17473 syscalls:sys_enter_seccomp:op: FILTER, flags: 0x2, uargs: 0x7fffc57e0e40
```

```
1.035 demo/17473 syscalls:sys_exit_seccomp:0x0
```

```
...
```

```
1.060 demo/17473 probe:seccomp_run_filters__return:(ffffffff8617e850 <- fffffffff8617ef4e)  
arg1=0x50001
```

```
1.090 demo/17473 syscalls:sys_exit_read:0xffffffffffffffff
```

```
...
```

# Perf record

```
$ sudo perf record -e 'probe:seccomp_run_filters,syscalls:*' demo ./prog.bpf
```

```
$ sudo perf script
```

# Demo - simple tracing with event trace

# Event trace

```
$ sudo perf trace -e 'tcp:*' demo http://ifconfig.co/json > /dev/null
```

```
$ sudo perf list |grep skb
```

```
$ sudo perf trace -e 'tcp:*,skb:*' demo http://ifconfig.co/json > /dev/null
```

```
$ sudo perf trace -e 'tcp:*,skb:*,udp:*' ./a.out http://ifconfig.co/json > /dev/null
```

# Event trace

- Generic event trace
- Manually defined trace points
- Manually invoked
- Helper macro `TRACE_EVENT(....)`
- Invoked as `trace_<name>()`
- Example:
  - `include/trace/events/skb.h`, `TRACE_EVENT(kfree_skb)`
  - `net/core/skbuff.c`, `trace_kfree_skb()`

# Demo - deeper tracing with uprobes & events

# uprobe

- Userspace probes
- Dynamically patched userspace code
- No modifications needed

```
$ perf probe -x <binary> <symbol> <format>
```

```
$ sudo perf probe -x /usr/lib/libsoup-2.4.so.1.8.0 soup_message_body_new
```

```
$ sudo perf probe -x /usr/lib/libsoup-2.4.so.1.8.0 soup_message_new
```

```
$ sudo perf probe -x /usr/lib/libgobject-2.0.so g_object_unref
```

```
$ sudo perf probe -x /usr/lib/libsoup-2.4.so.1.8.0 \
```

```
    'soup_message_body_new%return' '$retval:u32'
```



# uprobe - simple tracing

```
$ sudo perf trace \  
-e 'probe_libsoup:soup_session_send_message,probe_libsoup:soup_session_send_message__return' \  
./demo http://ifconfig.co/json > /dev/null
```

```
$ sudo perf trace -g -e ... ./demo http://ifconfig.co/json > /dev/null
```

# Demo - Userspace Statically Defined Trace

# USDT - Statically Defined Trace

```
#include <sys/sdt.h>
...
void foo(size_t cnt) {
    DTRACE_PROBE1(demo, foo, cnt);
    if (cnt % 2 == 0) printf("foo\n");
    else printf("oof\n");
}
void bar(size_t cnt) {
    DTRACE_PROBE1(demo, bar, cnt);
    if (cnt % 2 == 0) printf("bar\n");
    else printf("rab\n");
}
int main(int argc, char *argv[]) {
    int cnt = 0;
    while (1) {
        foo(cnt); bar(cnt); sleep(1); cnt++;
    }
    return 0;
}
```

# USDT

```
$ perf buildid-cache --add demo
```

```
$ perf list | grep sdt_demo
```

```
$ perf probe add sdt_demo:foo; perf probe add sdt_demo:bar
```

```
$ perf trace -e sdt_demo:foo,sdt_demo:bar ./demo
```

# Demo - eBPF + bcc + tracing

# Use eBPF & kprobe hook to figure out DNS server

- DNS, UDP port 53
- Limit to AF\_INET (IPv4)
- Packets are send with sendmsg(), with socket of SOCK\_DGRAM, AF\_INET

```
$ grep udp_send < /proc/kallsyms
0000000000000000 t udp_send_skb.isra.4
0000000000000000 T udp_sendmsg          ← USE THIS ONE
0000000000000000 T udp_sendpage
0000000000000000 t udp_sendmsg.cold.15
0000000000000000 r __ksymtab_udp_sendmsg
0000000000000000 r __kstrtab_udp_sendmsg
```

# udp\_sendmsg

```
int udp_sendmsg(struct sock *sk, struct msghdr *msg, size_t len)
{
    struct inet_sock *inet = inet_sk(sk);
    ...
    DECLARE_SOCKADDR(struct sockaddr_in *, usin, msg->msg_name);
    ...
    int connected = 0;
    __be32 daddr, faddr, saddr;
    __be16 dport;
    ...
    if (usin) {
        ...
        daddr = usin->sin_addr.s_addr;
        dport = usin->sin_port;
    } else {
        if (sk->sk_state != TCP_ESTABLISHED)
            return -EDESTADDRREQ;
        daddr = inet->inet_daddr;
        dport = inet->inet_dport;
        connected = 1;
    }
    ...
}
```

# bcc - BPF Compiler Collection

- Wrapper around calls to llvm
- Constrained C variant
- Hooks into tracing subsystem
- Receiving data from the tracing subsystem
- Processing of simple trace\_pipe
- C++, Python bindings



# bcc - kprobe

```
#include <uapi/linux/ptrace.h>
#include <net/sock.h>
#include <linux/socket.h>
#include <linux/in.h>
#include <linux/in6.h>

struct data_t {
    u32 family;
    u32 pid;
    u32 port;
    struct in_addr in;
};

// BPF table for pushing out data
BPF_PERF_OUTPUT(events);
```

# bcc - kprobe

```
/* UDP in IPv4 */
int kprobe__udp_sendmsg(struct pt_regs *ctx, struct sock *sk, struct msghdr *msg, size_t len) {
    struct data_t event = {};
    u32 tgid = bpf_get_current_pid_tgid() >> 32;
    event.pid = tgid;
    event.family = sk->__sk_common.skc_family;
    if (sk->sk_state == TCP_ESTABLISHED) {
        event.port = sk->__sk_common.skc_dport;
        event.in.s_addr = sk->__sk_common.skc_daddr;
    } else {
        struct sockaddr_in *addr = (struct sockaddr_in *)msg->msg_name;
        event.port = addr->sin_port;
        event.in = addr->sin_addr;
    }

    events.perf_submit(ctx, &event, sizeof(event));
    return 0;
}
```

Thank you!